Sarah Chen

AI/ML Engineer | sarah.chen@mleng.io

Professional Summary

AI/ML Engineer with 4+ years of experience training and deploying deep learning models, fine-tuning LLMs, and building low-latency RAG architectures in production.

Technical Skills & Technologies

Languages & Tools: Python, PyTorch, HuggingFace Transformers, LangChain, ChromaDB, FastAPI, Docker, CUDA, MLflow, Scikit-Learn

Domain Competencies: Deep Learning, Retrieval-Augmented Generation, NLP, Optimization, Model Quantization

Professional Experience

• Senior Machine Learning Engineer @ NeuralScale (2022 - Present): Architected an enterprise RAG pipeline serving 2M daily queries with sub-120ms p95 latency. Optimized cross-encoder re-ranking and implemented semantic caching in Redis.

• ML Engineer @ DataVanguard (2020 - 2022): Fine-tuned Llama and BERT models for classification and summarization. Mitigated severe training class imbalance using focal loss and SMOTE, boosting minority class F1-score from 0.58 to 0.83.

• Applied AI Researcher @ TechLabs (2019 - 2020): Researched gradient descent optimization variants (AdamW, Lion) and regularized deep CNNs using DropPath, Mixup, and BatchNorm.

Education

M.S. in Computer Science (Machine Learning Concentration), University of Washington